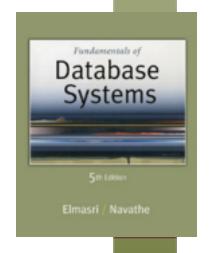


5th Edition

Elmasri / Navathe

Chapter 24

Enhanced Data Models for Advanced Applications





Outline

- Active database & triggers
- Temporal databases
- Spatial and Multimedia databases
- Deductive databases

Active Database Concepts and Triggers

Generalized Model for Active Databases and Oracle **Triggers**

- Triggers are executed when a specified condition occurs during insert/delete/update
 - Triggers are action that fire automatically based on these conditions

Event-Condition-Action (ECA) Model

- Triggers follow an Event-condition-action (ECA) model
 - Event:
 - Database modification
 - E.g., insert, delete, update),
 - Condition:
 - Any true/false expression

- Optional: If no condition is specified then condition is always true
- Action:
 - Sequence of SQL statements that will be automatically executed

Trigger Example

Generalized Model (contd.)

When a new employees is added to a department, modify the Total_sal of the Department to include the new employees salary

Condition

- Logically this means that we will CREATE a TRIGGER, let us call the trigger Total_sal1
 - This trigger will execute AFTER INSERT ON Employee table

- It will do the following FOR EACH ROW
 - WHEN NEW.Dno is NOT NULL
 - The trigger will UPDATE DEPARTMENT
 - By SETting the new Total_sal to be the sum of
 - •old Total_sal and NEW. Salary
 - WHERE the Dno matches the NEW.Dno;

Example: Trigger

Definition

DEPARTMENT

Can be CREATE or ALTER

CREATE TRIGGER Total_sal1 Can be FOR,
AFTER, INSTEAD

AFTER INSERT ON Employee of Can be INSERT,
UPDATE, DELETE The condition

WHEN (NEW.Dno is NOT NULL) UPDATE

SET Total_sal = Total_sal + NEW. Salary WHERE Dno = NEW.Dno;

The action

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24- 7

CREATE or ALTER TRIGGER

- CREATE TRIGGER <name>
 - Creates a trigger
- ALTER TRIGGER <name>
 - Alters a trigger (assuming one exists)
- CREATE OR ALTER TRIGGER <name>

- Creates a trigger if one does not exist
- Alters a trigger if one does exist
 - Works in both cases, whether a trigger exists or not

Conditions

- AFTER
 - Executes after the event
- BEFORE
 - Executes before the event
- INSTEAD OF
 - Executes instead of the event

- Note that event does not execute in this case
 - E.g., used for modifying views

Row-Level versus Statement-level

- Triggers can be
 - Row-level
 - FOR EACH ROW specifies a row-level trigger
 - Statement-level
 - Default (when FOR EACH ROW is not specified)
- Row level triggers
 - Executed separately for each affected row
- Statement-level triggers

Execute once for the SQL statement,

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24- 10

Condition

- Any true/false condition to control whether a trigger is activated on not
 - Absence of condition means that the trigger will always execute for the even
 - Otherwise, condition is evaluated
 - before the event for BEFORE trigger
 - after the event for AFTER trigger

Action

- Action can be
 - One SQL statement
 - A sequence of SQL statements enclosed between a BEGIN and an END
- Action specifies the relevant modifications

Triggers on Views

Generalized Model (contd.)

 INSTEAD OF triggers are used to process view modifications

Active Database Concepts and Triggers

Design and Implementation Issues for Active Databases

• An active database allows users to make the

following changes to triggers (rules)

- Activate
- Deactivate
- Drop

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24- 14

Active Database Concepts and Triggers

Design and Implementation Issues for Active Databases

An event can be considered in 3 ways

- Immediate consideration
- Deferred consideration
- Detached consideration

Active Database Concepts and Triggers

Design and Implementation Issues (contd.)

- Immediate consideration
 - Part of the same transaction and can be one of the following depending on the situation

- Before
- After
- Instead of
- Deferred consideration
 - Condition is evaluated at the end of the transaction
- Detached consideration
 - Condition is evaluated in a separate transaction

Active Database Concepts and Triggers

Potential Applications for Active

- Databases Notification
 - Automatic notification when certain condition occurs
 - Enforcing integrity constraints

- Triggers are smarter and more powerful than constraints
- Maintenance of derived data
 - Automatically update derived data and avoid anomalies due to redundancy
 - E.g., trigger to update the Total_sal in the earlier example

Active Database Concepts and Triggers

Triggers in SQL-99

Can alias variables inside the REFERENCINFG clause

Active Database Concepts and Triggers

Trigger examples

```
T1: CREATE TRIGGER Total sal1
    AFTER UPDATE OF Salary ON EMPLOYEE
    REFERENCING OLD ROW AS O, NEW ROW AS N
    FOR EACH ROW
    WHEN ( N.Dno IS NOT NULL )
    UPDATE DEPARTMENT
    SET Total_sal = Total_sal + N.salary - O.salary
    WHERE Dno = N.Dno;
T2: CREATE TRIGGER Total sal2
    AFTER UPDATE OF Salary ON EMPLOYEE
    REFERENCING OLD TABLE AS O, NEW TABLE AS N
    FOR EACH STATEMENT
    WHEN EXISTS ( SELECT * FROM N WHERE N.Dno IS NOT NULL ) OR
            EXISTS ( SELECT * FROM O WHERE O.Dno IS NOT NULL )
    UPDATE DEPARTMENT AS D
    SET D.Total_sal = D.Total_sal
    + ( SELECT SUM (N.Salary) FROM N WHERE D.Dno=N.Dno )
    - ( SELECT SUM (O.Salary) FROM O WHERE D.Dno=O.Dno )
    WHERE Dno IN ((SELECT Dno FROM N) UNION (SELECT Dno FROM O));
                                                                      Copyright ©
```

 $_{\rm 2007\;Ramez\;Elmasri\;and\;Shamkant\;B.\;Navathe}$ Slide 24- 19

Temporal Database Concepts

Time Representation, Calendars, and Time Dimensions Time is considered **ordered sequence** of **points** in some

granularity

 Use the term choronon instead of point to describe minimum granularity

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24- 20

Temporal Database Concepts

Time Representation, ... (contd.)

A calendar organizes time into different time

units for convenience.

- Accommodates various calendars
 - Gregorian (western)
 - Chinese
 - Islamic
 - Hindu
 - Jewish
 - Etc.

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24-21

Temporal Database Concepts

Time Representation, ... (contd.)

- Point events
 - Single time point event

- E.g., bank deposit
- Series of point events can form a time series data
- Duration events
 - Associated with specific time period
 - Time period is represented by start time and end time

Temporal Database Concepts

Time Representation, ... (contd.)

- Transaction time
 - The time when the information from a certain transaction becomes valid
- Bitemporal database

Databases dealing with two time dimensions

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24-23

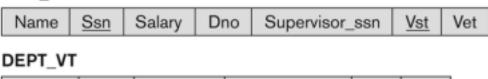
Temporal Database Concepts

Incorporating Time in Relational Databases Using Tuple Versioning

- Add to every tuple
 - Valid start time
 - Valid end time

Temporal Database Concepts





Manager_ssn

Total_sal

Total sal

Figure 24.7

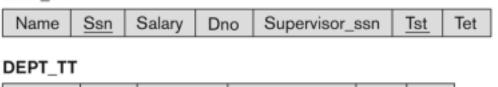
Different types of temporal relational databases. (a) Valid time database schema. (b) Transaction time database schema. (c) Bitemporal database schema.

(b) EMP_TT

Dname

Dno

Dno



(c) EMP BT



Manager_ssn

Vet

Tet

Vst

Tst

DEPT_BT

Dname



Copyright ©

 $_{\rm 2007\;Ramez\;Elmasri\;and\;Shamkant\;B.\;Navathe}Slide~24-~25$

Temporal Database Concepts

Figure 24.8

Some tuple versions in the valid time relations EMP_VT and DEPT_VT.

EMP_VT

Name	<u>Ssn</u>	Salary	Dno	Supervisor_ssn	<u>Vst</u>	Vet
Smith	123456789	25000	5	333445555	2002-06-15	2003-05-31
Smith	123456789	30000	5	333445555	2003-06-01	Now
Wong	333445555	25000	4	999887777	1999-08-20	2001-01-31
Wong	333445555	30000	5	999887777	2001-02-01	2002-03-31
Wong	333445555	40000	5	888665555	2002-04-01	Now
Brown	222447777	28000	4	999887777	2001-05-01	2002-08-10
Narayan	666884444	38000	5	333445555	2003-08-01	Now

. . .

DEPT_VT

Dname	<u>Dno</u>	Manager_ssn	<u>Vst</u>	Vet
Research	5	888665555	2001-09-20	2002-03-31
Research	5	333445555	2002-04-01	Now

 ${}_{\text{Copyright @ 2007 Ramez Elmasri and Shamkant B. Navathe}}Slide~24-~26$

Temporal Database Concepts

Incorporating Time in Object-Oriented Databases Using Attribute Versioning

- A single complex object stores all temporal changes of the object
- Time varying attribute
 - An attribute that changes over time
 - E.g., age
- Non-Time varying attribute
 - An attribute that does not changes over time
 - E.g., date of birth

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24-27

Spatial and Multimedia Databases

Spatial Databases

Multimedia Database Concepts

- Keep track of objects in a multi-dimensional space
 - Maps
 - Geographical Information Systems (GIS)
 - Weather
- In general spatial databases are n-dimensional
 - This discussion is limited to 2-dimensional spatial databases

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24-29

- Typical Spatial Queries
 - Range query: Finds objects of a particular type within a particular distance from a given location
 - E.g., Taco Bells in Pleasanton, CA
 - Nearest Neighbor query: Finds objects of a particular type that is nearest to a given location
 - E.g., Nearest Taco Bell from an address in Pleasanton, CA **Spatial joins** or overlays: Joins objects of two types based on some spatial condition (intersecting, overlapping, within certain distance, etc.)
 - E.g., All Taco Bells within 2 miles from I-680.

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24-30

Spatial Databases

R-trees

- Technique for typical spatial queries
- Group objects close in spatial proximity on the same leaf nodes of a tree structured index
- Internal nodes define areas (rectangles) that cover all areas of the rectangles in its subtree.

Quad trees

Divide subspaces into equally sized areas

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24-31

Multimedia Databases

Multimedia Database Concepts

- In the years ahead multimedia information systems are expected to dominate our daily lives.
 - Our houses will be wired for bandwidth to handle interactive multimedia. applications.
 - Our high-definition TV/computer workstations will have access to a large number of databases, including digital libraries, image and video databases that will distribute vast amounts of multisource multimedia content.

- Types of multimedia data are available in current systems
 - Text: May be formatted or unformatted. For ease of parsing structured documents, standards like SGML and variations such as HTML are being used.
 - Graphics: Examples include drawings and illustrations that are encoded using some descriptive standards (e.g. CGM, PICT, postscript).

- Types of multimedia data are available in current systems (contd.)
 - Images: Includes drawings, photographs, and so forth, encoded in standard formats such as bitmap, JPEG, and MPEG. Compression is built into JPEG and MPEG.
 - These images are not subdivided into components. Hence querying them by content (e.g., find all images containing circles) is nontrivial.
 - Animations: Temporal sequences of image or graphic data.

- Types of multimedia data are available in current systems (contd.)
 - Video: A set of temporally sequenced photographic data for presentation at specified rates—for example, 30 frames per second.
 - Structured audio: A sequence of audio components comprising note, tone, duration, and so forth.

- Types of multimedia data are available in current systems (contd.)
 - Audio: Sample data generated from aural recordings in a string of bits in digitized form. Analog recordings are typically converted into digital form before storage.

- Types of multimedia data are available in current systems (contd.)
 - Composite or mixed multimedia data: A combination of multimedia data types such as audio and video which may be physically mixed to yield a new storage format or logically mixed while retaining original types and formats. Composite data also contains additional control information describing how the information should be rendered.

Nature of Multimedia Applications:

- Multimedia data may be stored, delivered, and utilized in many different ways.
- Applications may be categorized based on their data management characteristics.

- Overview of Deductive Databases
- Prolog/Datalog Notation
- Datalog Notation
- Clausal Form and Horn Clauses
- Interpretation of Rules
- Datalog Programs and Their Safety
- Use the Relational Operations
- Evaluation of Non-recursive Datalog Queries

- Declarative Language
 - Language to specify rules
- Inference Engine (Deduction Machine)
 - Can deduce new facts by interpreting the rules
 - Related to logic programming
 - Prolog language (Prolog => Programming in logic)
 - Uses backward chaining to evaluate
 - Top-down application of the rules

- Speciation consists of:
 - Facts
 - Similar to relation specification without the necessity of including attribute names
 - Rules
 - Similar to relational views (virtual relations that are not stored)

- **Predicate** has
 - a name
 - a fixed number of arguments
 - Convention:
 - Constants are numeric or character strings
 - Variables start with upper case letters
 - E.g., SUPERVISE(Supervisor, Supervisee)
 - States that Supervisor SUPERVISE(s) Supervisee

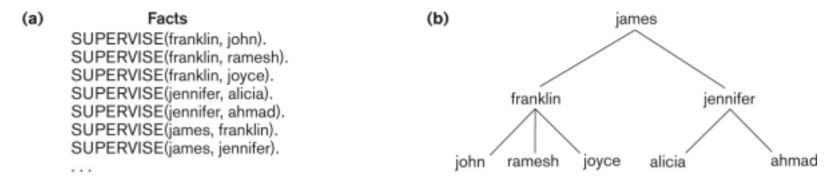
Rule

- Is of the form head :- body
 - where :- is read as if and only iff
 - E.g., SUPERIOR(X,Y) :- SUPERVISE(X,Y) ●
 - E.g., SUBORDINATE(Y,X) :- SUPERVISE(X,Y)

Query

- Involves a predicate symbol followed by y some variable arguments to answer the question
 - where :- is read as if and only iff
- E.g., SUPERIOR(james,Y)?
- E.g., SUBORDINATE(james,X)?

(a) Prolog notation (b) Supervisory tree



Rules

SUPERIOR(X, Y) := SUPERVISE(X, Y). SUPERIOR(X, Y) := SUPERVISE(X, Z), SUPERIOR(Z, Y). SUBORDINATE(X, Y) := SUPERIOR(Y, X).

Queries

SUPERIOR(james, Y)? SUPERIOR(james, joyce)?

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24-45

Datalog Notation

Datalog notation

- Program is built from atomic formulae
 - Literals of the form p(a1, a2, ... an) where
 - p predicate name
 - n is the number of arguments
 - Built-in predicates are included
 - •E.g., <, <=, etc.
- A literal is either
 - An atomic formula
 - •An atomic formula preceded by not

Clausal Form and Horn Clauses

- A formula can have quantifiers
 - Universal

Existential

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24-47

Clausal Form and Horn Clauses

 In clausal form, a formula must be transformed into another formula with the following characteristics

- All variables are universally quantified
- Formula is made of a number of clauses where each clause is made up of literals connected by logical ORs only
- Clauses themselves are connected by logical ANDs only.

Clausal Form and Horn Clauses

- Any formula can be converted into a clausal form
 - A specialized case of clausal form are horn clauses that can contain no more than one positive literal
- Datalog program are made up of horn clauses

Interpretation of Rules

- There are two main alternatives for interpreting rules:
 - Proof-theoretic
 - Model-theoretic

Interpretation of Rules

- Proof-theoretic
 - Facts and rules are axioms
 - Ground axioms contain no variables
 - Rules are deductive axioms
 - Deductive axioms can be used to construct new facts from existing facts

This process is known as theorem proving

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24-51

Proving a new fact

• Figure 24.12

```
    SUPERIOR(X, Y): - SUPERVISE(X, Y). (rule 1)
    SUPERIOR(X, Y): - SUPERVISE(X, Z), SUPERIOR(Z, Y). (rule 2)
    SUPERVISE(jennifer, ahmad). (ground axiom, given)
    SUPERVISE(james, jennifer). (ground axiom, given)
    SUPERIOR(jennifer, ahmad). (apply rule 1 on 3)
    SUPERIOR(james, ahmad). (apply rule 2 on 4 and 5)
```

 ${}_{\text{Copyright}\, @\, 2007\, \text{Ramez Elmasri and Shamkant B. Navathe}} Slide\,\, 24\text{--}\,\, 52$

Interpretation of Rules

- Model-theoretic
 - Given a finite or infinite domain of constant values, we assign the predicate every combination of values as arguments
 - If this is done fro every predicated, it is called interpretation

Interpretation of Rules

- Model
 - An interpretation for a specific set of rules
- Model-theoretic proofs
 - Whenever a particular substitution to the variables in the rules is applied, if all the predicated are true under the interpretation, the predicate at the

head of the rule must also be true

Minimal model

 Cannot change any fact from true to false and still get a model for these rules

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24-54

Minimal model Figure 24.13

Rules

SUPERIOR(X, Y) := SUPERVISE(X, Y).

SUPERIOR(X, Y) := SUPERVISE(X, Z), SUPERIOR(Z, Y).

Interpretation

Known Facts:

SUPERVISE(franklin, john) is true.

SUPERVISE(franklin, ramesh) is true.

SUPERVISE(franklin, joyce) is true.

SUPERVISE(jennifer, alicia) is true.

SUPERVISE(jennifer, ahmad) is true.

SUPERVISE(james, franklin) is true.

SUPERVISE(james, jennifer) is true.

SUPERVISE(X, Y) is false for all other possible (X, Y) combinations

Derived Facts:

SUPERIOR(franklin, john) is true.

SUPERIOR(franklin, ramesh) is true.

SUPERIOR(franklin, joyce) is true.

SUPERIOR(jennifer, alicia) is true.

SUPERIOR(jennifer, ahmad) is true.

SUPERIOR(james, franklin) is true.

SUPERIOR(james, jennifer) is true.

SUPERIOR(james, john) is true.

SUPERIOR(james, ramesh) is true.

SUPERIOR(james, joyce) is true.

SUPERIOR(james, alicia) is true.

SUPERIOR(james, ahmad) is true.

SUPERIOR(X, Y) is **false** for all other possible (X, Y) combinations

Copyright © 2007 Ramez Elmasri and Shamkant B. Navathe Slide 24- 55

Datalog Programs and Their Safety

- Two main methods of defining truth values
 - Fact-defined predicates (or relations)
 - Listing all combination of values that make a predicate true
 - Rule-defined predicates (or views)
 - Head (LHS) of 1 or more Datalog rules, for example
 Figure 24.15

```
SUPERIOR(X, Y) := SUPERVISE(X, Y).

SUPERIOR(X, Y) := SUPERVISE(X, Z), SUPERIOR(X, Y).

SUBORDINATE(X, Y) := SUPERIOR(Y, X).

SUPERVISOR(X) := EMPLOYEE(X), SUPERVISE(X, Y).

OVER_40K_EMP(X) := EMPLOYEE(X), SALARY(X, Y), Y> = 40000.

UNDER_40K_SUPERVISOR(X) := SUPERVISOR(X), NOT(OVER_40_K_EMP(X)).

MAIN_PRODUCTX_EMP(X) := EMPLOYEE(X), WORKS_ON(X, Y), Y> = 20.

PRESIDENT(X) := EMPLOYEE(X), NOT(SUPERVISE(Y, X)).
```

Datalog Programs and Their Safety

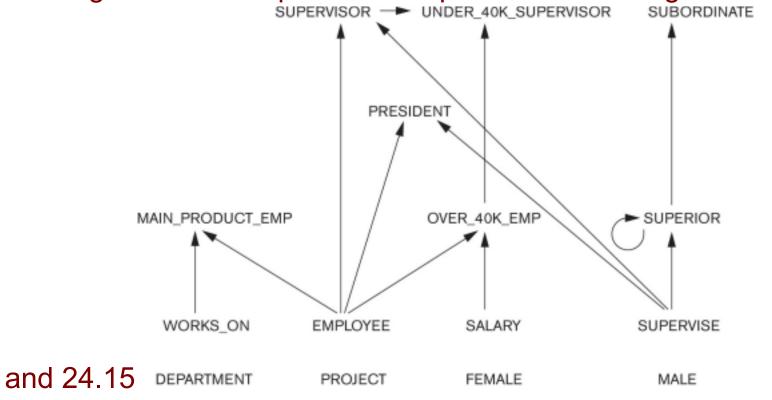
- A program is safe if it generates a finite set of facts
 - Fact-defined predicates (or relations)
 - Listing all combination of values that make a predicate true
 - Rule-defined predicates (or views)
 - Head (LHS) of 1 or more Datalog rules, for example Figure 24.15

- Many operations of relational algebra can be defined in the for of Datalog rules that defined the result of applying these operations on database relations (fact predicates)
 - Relational queries and views can be easily specified in Datalog

Evaluation of Non-recursive Datalog Queries

 Define an inference mechanism based on relational database query processing concepts

See Figure 24.17 on predicate dependencies for Figs 24.14





- Active database & triggers
- Temporal databases
- Spatial and Multimedia databases
- Deductive databases